

# SVM-vs-Knn

Giulio Bagnoli

November 2019

**Abstract:** The purpose of this document is to show how classification work through the use of KNN and SVM and how the results obtained change with the tuning of the various hyperparameters of the algorithms. To do that *sklearn.datasets.load\_wine*[13] was used and some tests are been conducted:

- For  $k \in [1, 3, 5, 7]$  KNN was built.
- For  $C \in [0.001, 0.01, 0.1, 1, 10, 100, 1000]$  linear SVM was built.
- For  $C \in [0.001, 0.01, 0.1, 1, 10, 100, 1000]$  SVM with *RBF* kernel was built.
- A grid search was performed for  $\gamma$  and  $C$  for each possible pair an SVM with *RBF* kernel was built.
- A grid search was performed for  $\gamma$  and  $C$  but this time perform 5-fold validation.

For each experiment the best model was used to evaluate the test set.

## 1 Classification

In machine learning and statistics[9], classification is the problem of identifying to which of a set of categories a new observation belongs, on the basis of a training set of data containing observations (or instances) whose category membership is known. Examples are assigning a given email to the "spam" or "non-spam" class, and assigning a diagnosis to a given patient based on observed characteristics of the patient (sex, blood pressure, presence or absence of certain symptoms, etc.). Classification is an example of pattern recognition.

In the terminology of machine learning, classification is considered an instance of supervised learning. A supervised learning algorithm analyses the

training data and produces an inferred function, which can be used for mapping new examples. An optimal scenario will allow for the algorithm to correctly determine the class labels for unseen instances. This requires the learning algorithm to generalize from the training data to unseen situations in a "reasonable" way.

An algorithm that implements classification is known as a classifier. The term "classifier" sometimes also refers to the mathematical function, implemented by a classification algorithm, that maps input data to a category.

The classification can be considered as two separate problems: the binary classification and the multi-class classification. In binary classification, a better understood activity, only two classes are involved, while the multi-class classification involves the assignment of an object to one of the different classes. There are algorithms that can solve only the first problem, such as the Perceptron.

A wide range of supervised learning algorithms are available, each with its strengths and weaknesses. There is no single learning algorithm that works best on all supervised learning problems.

## 2 K-Nearest Neighbour (KNN)

Neighbors-based classification is a type of instance-based learning or non-generalizing learning: it does not attempt to construct a general internal model, but simply stores instances of the training data [4]. In pattern recognition, the KNN algorithm is used for classification and regression. In the classification process, the output is a class membership.

Classification is done by a majority vote of  $K$  nearest neighbours. The basic nearest neighbors classification uses uniform weights, under some circumstances, it is better to weight the neighbors

such that nearer neighbors contribute more to the vote. Only uniform weights were used during the homework.

$K$  in KNN is a hyperparameter that must be chosen when the model is built. It can be interpreted as a control variable for the prediction model. Each dataset has its own requirements but usually the value of  $K$  is odd when the number of classes is two, in order to avoid the same quantity of votes for both classes. In the case of a small number of neighbors, the noise will have a higher influence on the result; on the other hand, a large number of neighbors makes the classification boundaries less distinct and the new data are influenced, especially in a sparse space, by points distant from them.




Figure 1: How KNN classifies new data.

To define which are the nearest  $K$  neighbors it can be used different distance metrics like:

- Minkowski  $L_p(x, y) = \left( \sum_{i=1}^d |x_i - y_i|^p \right)^{\frac{1}{p}}$
- Manhattan  $L_1(x, y) = \left( \sum_{i=1}^d |x_i - y_i| \right)$
- Euclidean  $L_2(x, y) = \left( \sum_{i=1}^d (x_i - y_i)^2 \right)^{\frac{1}{2}}$
- Chebyshev  $L_\infty(x, y) = \left( \max_i |x_i - y_i| \right)$

It is possible to choose different metrics to measure the distance between the various data because there was no metric of the optimal distance that could be used for all types of datasets as the free lunch theorem states.[5]

There is a possibility that using different distance metrics we might get better results. So, in non-

probabilistic algorithm like KNN distance metrics plays an important role.

Nearest Neighbors Classifier is attractive for classification because it can naturally handle multi-class problems without any increase in the model size, and does not introduce additional parameters that require fine-tuning by the user.

### 3 Support Vector Machine (SVM)

A Support Vector Machine is a discriminative classifier formally defined by a separating hyperplane. In other words, given labeled training data, the algorithm outputs an optimal hyperplane which categorizes new examples. Intuitively, the best separation is achieved by the hyperplane that has the largest distance to the nearest training point of any class, since in general the larger this distance, the smaller the generalization error of the classifier [2].

If the training data is linearly separable, we can select two parallel hyperplanes that separate the two classes of data. The region bounded by these two hyperplanes is called the "margin", and the maximum-margin hyperplane is the hyperplane that lies halfway between them.




Figure 2: Hyperplane in two dimensional space.

In binary classification SVM has as inputs a training dataset of  $n$  points of the form:

$$(x_1, y_1), \dots, (x_n, y_n) \text{ with } x \in R^n, y \in \{-1, 1\} \quad (1)$$

any hyperplane can be written as the set of points  $x$  satisfying:

$$w \cdot x - b = 0 \quad (2)$$

where  $w$  is the normal vector to the hyperplane. The distance between the hyperplane that delimit the margin is

$$\frac{2}{\|w\|} \quad (3)$$

so to maximize the margin width we want to minimize  $w$ . We also have to prevent data points from falling into the margin, to do that we add the following constraint:

$$w \cdot x_i - b \geq 1 \quad \text{if } y_i = 1 \quad (4)$$

$$w \cdot x_i - b \leq -1 \quad \text{if } y_i = -1 \quad (5)$$

(4) and (5) can be put together as:

$$y_i(w \cdot x_i - b) \geq 1 \quad (6)$$

Combining the optimization problem with the constraint above we get:

$$\begin{cases} \min \|w\| \\ y_i(w \cdot x_i - b) \geq 1 \end{cases} \quad (7)$$

The maximum-margin hyperplane is completely determined by  $x_i$  that lie nearest to it. These  $x_i$  are called *support vectors*.

In real applications there is not always a margin, therefore classes are not always linearly separable in the space of features through a hyperplane. Positive slack variables  $\xi_i$  are introduced to relax the margin constraint. The problem becomes:

$$\begin{cases} \min \frac{1}{2} \|w\|^2 + C \sum \xi_i \\ y_i(w \cdot x_i - b) \geq 1 - \xi_i \end{cases} \quad (8)$$

Here,  $C$  is a hyperparameter that decides the trade-off between maximizing the margin and minimizing the mistakes. When  $C$  is small, to the classification mistakes are given less importance and focus is more on maximizing the margin, whereas when  $C$  is large, the focus is more on avoiding misclassification at the expense of keeping the margin small.

## 4 Description of the project

The aim of the work is to see the different behaviors between KNN and SVM and how they change as hyperparameters change.

To do this, the first two features were selected from the **wine** dataset, the data were randomly divided into three sets:

- *Train*: the sample of data used to fit the model. The classifier sees and learns from this data.
- *Validation*: the sample of data used to provide an evaluation of a model fit on the training dataset while tuning model hyperparameters. So the validation set in a way affects a model, but indirectly.
- *Test*: The sample of data used to provide an unbiased evaluation of a final model fit on the training dataset [12]. It is only used when the model is completely trained.

The first set is composed by 50% of the data, the second one by 20% and the last one by 30%.

Since the range of values of raw data may vary widely, in some machine learning algorithms, objective functions will not work properly without normalization. For example, KNN calculates the distance between the new point and the elements contained in train set so if one of the features has a broad range of values, the distance will be governed by this particular feature. Therefore, the range of all features should be normalized so that each feature contributes approximately proportionately to the final distance. To do this the class `sklearn.preprocessing.StandardScaler` [8] was used which normalizes the data as follows:

$$z = \frac{(x - u)}{s} \quad (9)$$

where  $u$  is the mean of the union of the training and validation samples and  $s$  is the standard deviation of the same data.

For  $k \in [1, 3, 5, 7]$  the model was built using the `sklearn.neighbors.KNeighborsClassifier` [3] to which the train set was passed. Euclidean distance was used as metric. Through the use of the model the decision boundaries were printed and the validation set was evaluated.




Figure 3: Decision boundaries when  $K=1$ .

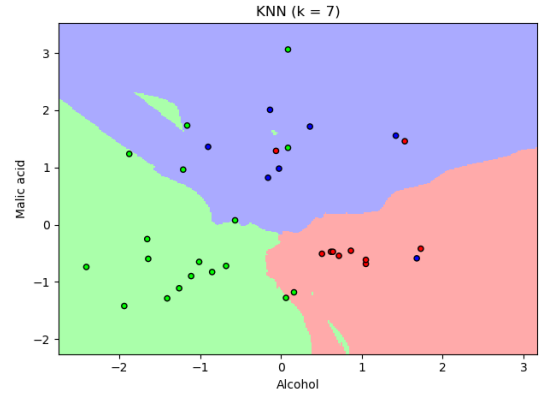


Figure 6: Decision boundaries when  $K=7$ .




Figure 4: Decision boundaries when  $K=3$ .




Figure 5: Decision boundaries when  $K=5$ .

In the images from 3 to 6 it is visible how KNN classifies the elements of the validation set based on the decision boundaries. Decision boundaries change because as  $K$  increases, they are influenced by more distant elements. In particular it is visible that for small  $K$  they are more jagged whereas for larger  $K$  they tend to be smoother.




Figure 7: Accuracy on validation set for different values of  $K$ .

Figure 7 shows that, for this data distribution, the classifier for  $K = 3$  and  $K = 5$  obtains the worst result (71.42%) on the validation set, whereas the accuracy improves considerably for  $K = 1$ . The accuracy for this value of the model is (80.007%). The classifier, trained with  $k = 1$  and evaluated on the test set has an accuracy of 77.78%. It is very important to note that the results thus obtained depend on the random shuffle

performed at the beginning, and that, with other shuffles, different results are obtained.

Later for  $c \in [0.001, 0.01, 0.1, 1, 10, 100, 1000]$  the model was built using the class `sklearn.svm.SVC`[11] and the linear kernel. `sklearn.svm.SVC` implement the “one-against-one” approach for multi- class classification. If  $n$  is the number of classes, then

$$\frac{n \cdot (n - 1)}{2} \quad (10)$$

classifiers are constructed and each one trains data from two classes. As above the model was trained with the train set and decision boundaries were printed.

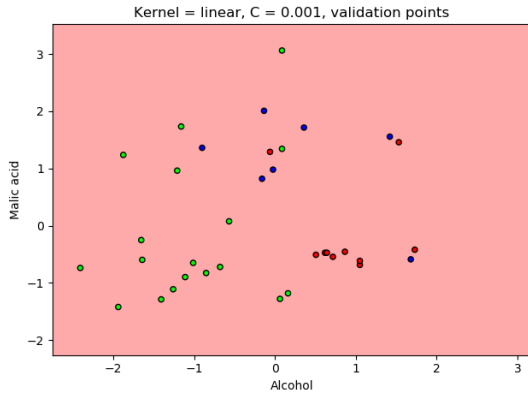


Figure 8: Decision boundaries when  $C=0.001$ .




Figure 9: Decision boundaries when  $C=0.01$ .

The decision boundaries obtained for  $C = 0.001$  is due to the fact that the weight given to the errors committed in the train set,  $C$ , is so low that it does not affect the objective function of the SVM. The only purpose of the classifier is therefore to obtain wider margins that are possible.




Figure 10: Decision boundaries when  $C=0.1$ .




Figure 11: Decision boundaries when  $C=1$ .




Figure 12: Decision boundaries when  $C=10$ .

On the other hand, when it was assign larger values to  $C$  the model tries to make fewer mistakes but reduces the margins. For  $C > 10$  the diagrams are not shown because there are no changes. In all cases, since the linear kernel was used the hyper-planes are formed by straight lines.




Figure 13: Accuracy on validation set for different values of  $C$ .

The figure 13 shows the evolution of the accuracy of the model evaluated on the validation set when  $C$  changes. Accuracy reaches its maximum value (77.14%) for  $C = 0.1$  and for  $C = 1$ . For values higher than  $C$  the accuracy of the model decreases because the classifier gives too much weight to the mistakes made during the train and therefore to reduce these errors reduces the margins. However, this makes it easier for SVM to make mistakes by classifying data it has never

seen. The classifier, trained with  $c = 0.1$ , and evaluated on the test set has an accuracy of 87.04%.

Since it is not possible to find a linear decision boundary for the classification problem then the data were projected in to a higher dimension from the original space to try to find a hyperplane in the new dimension that helps to classify the data. To do this the SVM kernel has been changed, from linear to RBF. So the dot product in the linear classifier is replace by:

$$K(x, x') = \exp \left( -\frac{\|x - x'\|^2}{2\gamma^2} \right) \quad (11)$$

$\gamma$  parameter can be seen as the inverse of the radius of influence of samples selected by the model as support vectors. All tests performed on the SVM have been repeated leaving the default value for  $\gamma$ :

$$\gamma = \frac{1}{(n\_features * \sigma)} \quad (12)$$

where  $\sigma$  is the variance of the train set.




Figure 14: Decision boundaries when kernel is RBF and  $C=0.1$ .

The decision boundary calculated with  $C = 0.001$  and the RBF kernel is the same as that calculated with the same  $C$  and linear kernel, because in both cases too little weight has been given to the misclassification.




Figure 15: Decision boundaries when kernel is RBF and  $C = 0.1$ .




Figure 17: Decision boundaries when kernel is RBF and  $C=1$ .




Figure 16: Decision boundaries when kernel is RBF and  $C=1000$ .

The maximum accuracy, calculated on the validation set, for the various values of  $C$ , improves passing from 77,14%, for the SVM with linear kernel, to 80.00% always for  $C = 0.1$ . The accuracy on the test set of the classifier, trained with the best  $C$  on the validation set, increases from 75.92% to 79.63%. This is due to the fact that the data is better separable in a higher dimension.

To try to further improve the results both  $\gamma$  and  $C$  are tuned at the same time. The set of values tested for  $\gamma$  is  $[0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1, 10]$  and for  $C$  is  $[0.01, 0.05, 0.1, 0.5, 1, 5, 10, 100, 1000]$ . The lower value of  $C$  was not taken, as the previous experiments show that with this value a worse model was obtained compared to the other values.

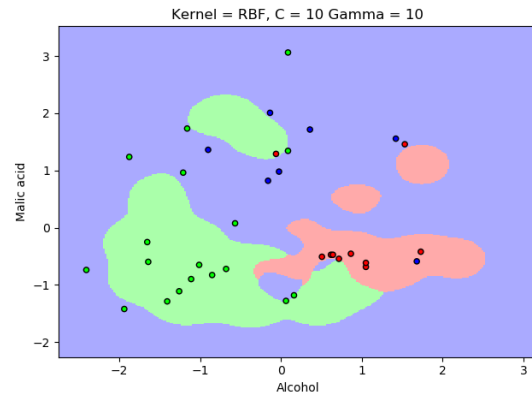


Figure 18: Decision boundaries when kernel is RBF and  $C=10$  and  $\gamma = 10$ .

The graphs, for a sufficiently high  $C$ , show how the decision boundaries are no longer formed by straight lines but by curve lines, this is due to the use of a non-linear kernel.

The behavior of the model is very sensitive to the  $\gamma$  parameter. If  $\gamma$  is too large, the radius of the area of influence of the support vectors only includes the support vector itself and no amount of regularization with  $C$  will be able to prevent overfitting.[6]

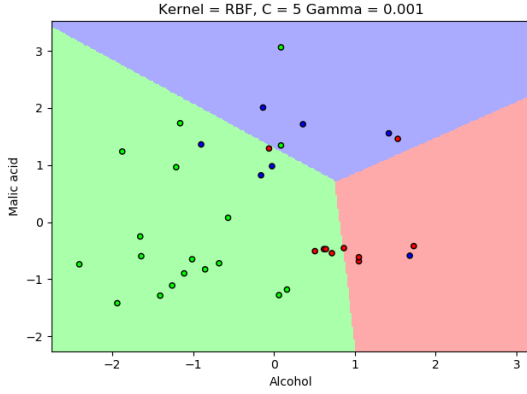


Figure 19: Decision boundaries when kernel is RBF and  $C=5$  and  $\gamma = 0.001$ .

When  $\gamma$  is very small, the model is too constrained and cannot capture the complexity or “shape” of the data. The region of influence of any selected support vector would include the whole training set. The resulting model will behave similarly to a linear model.




Figure 20: Decision boundaries when kernel is RBF and  $C=1000$  and  $\gamma = 0.05$ .

The couple of hyperparameters that gives a better accuracy (85.71%) on the validation set is ( $C = 1000, \gamma = 0.05$ ). The accuracy on the test set of

the classifier, trained with these value, is 81.48%. The model is more precise than the previous one because both  $\gamma$  and  $C$  have been tuned at the same time.

After merge of training and validation split it was done a grid search for  $\gamma$  and  $C$  perform 5-fold validation. To do k-fold validation was used the function `sklearn.model_selection.cross_val_score[1]`.

In k-fold cross-validation, the train set is randomly partitioned into  $k$  equal size subsamples. Of the  $k$  subsamples, a single subsample is retained as the validation data for testing the model, and the remaining  $k - 1$  subsamples are used as training data. The cross-validation process is then repeated  $k$  times, with each of the  $k$  subsamples used exactly once as the validation data. The average of the results from the folds was calculated to produce a single estimation of model. The advantage of this method is that all observations are used for both training and validation, and each observation is used for validation exactly once. Furthermore when k-fold cross-validation is used it creates  $k$  different models and test them on  $k$  different validation sets, it makes the performance of the algorithm more stable. Indeed when only one single appraisal is done on validation set, this result may be because of chance or a biased test set for some reason. By training  $k$  different models is more understandable what is happening.

The couple of hyperparameters that gives a better accuracy (79.05%) on the validation set is ( $C = 100, \gamma = 0.05$ ). The accuracy on the test set of the classifier, trained with these value, is 87.04%. Accuracy increases on the test set, this could be due to the fact that by trying different combinations of data the model has learned of patterns not previously known.

Since irrelevant data components can adversely affect the model's performance, it was decided to use a feature selection technique to determine a feature ranking and use the best and the worst couple to retrain both models, KNN and SVM. To create the ranking, which is independent of the classifier, the ReliefF[7] algorithm was used, this is a filter type method that assigns score to the features. This score goes between -1 and 1, where a greater weight indicates a greater relevance of this feature. The score was calculated using only train and validation set. Below are showed the



scores for each features.

1. od280/od315 of diluted wines 0.190
2. flavanoids 0.176
3. proline 0.171
4. alcohol 0.152
5. color intensity 0.133
6. total phenols 0.114
7. hue 0.110
8. malic acid 0.081
9. nonflavanoid phenols 0.061
10. alcalinity of ash 0.054
11. proanthocyanins 0.047
12. magnesium 0.030

All the experiments performed on the features (*alcohol, malic acid*) have been repeated for the couples of features (*proanthocyanins, magnesium*) and (*od280/od315 of diluted wine, flavanoids*).

	<i>proanthocyanins, magnesium</i>	<i>od280..., flavanoids</i>
KNN ( $K$ )	57.40%	81.48%
linear SVM ( $C$ )	50.00%	83.33%
rbf SVM( $C$ )	46.29%	85.18%
rbf SVM( $C, \gamma$ )	55.55%	85.18%
Cross-fold( $C, \gamma$ )	55.55%	85.18%

Table 1: Accuracy obtained on the test set.

The accuracy obtained on the test set are reported in the tab 1: it can be seen how the use of different features strongly changes the results obtained. As expected the features with higher scores had better results than the other two pairs tested.

## References

- [1] *cross\_val\_score*. URL: [https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.cross\\_val\\_score.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.cross_val_score.html).
- [2] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. New York: Springer, 2008.
- [3] *KNeighborsClassifier*. URL: <https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html#sklearn.neighbors.KNeighborsClassifier>.
- [4] *Nearest Neighbors*. URL: <https://scikit-learn.org/stable/modules/neighbors.html#classification>.
- [5] V. B. Surya Prasatha et al. “Effects of Distance Measure Choice on KNN Classifier Performance”. In: (). URL: <https://arxiv.org/pdf/1708.04321.pdf>.
- [6] *RBF SVM parameters*. URL: [https://scikit-learn.org/stable/auto\\_examples/svm/plot\\_rbf\\_parameters.html](https://scikit-learn.org/stable/auto_examples/svm/plot_rbf_parameters.html).
- [7] *ReliefF*. URL: <https://epistasislab.github.io/scikit-rebate/using/>.
- [8] *StandardScaler*. URL: <https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.StandardScaler.html>.
- [9] *Statistical classification*. URL: [https://en.wikipedia.org/wiki/Statistical\\_classification](https://en.wikipedia.org/wiki/Statistical_classification).
- [10] *Supervised Learning*. URL: [https://en.wikipedia.org/wiki/Supervised\\_learning](https://en.wikipedia.org/wiki/Supervised_learning).
- [11] *SVM*. URL: <https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html#sklearn.svm.SVC>.
- [12] Tatiana Tommasi. *3aiml-lecture*. P. 45.
- [13] *Wine dataset*. URL: [https://scikit-learn.org/stable/modules/generated/sklearn.datasets.load\\_wine.html](https://scikit-learn.org/stable/modules/generated/sklearn.datasets.load_wine.html).